

✓ Multiclass Text Classification with

Logistic Regression Implemented with PyTorch and CE Loss

First, we will do some initialization.

```
from google.colab import drive
drive.mount('/content/drive')
!pwd
```

```
↗ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True)
```

```
%cd "drive"
%cd "MyDrive"
%cd "ag_new_csv"
!ls
```

```
↗ /content/drive
/content/drive/MyDrive
/content/drive/MyDrive/ag_new_csv
classes.txt test.csv train.csv
```

```
import random
import torch
import numpy as np
import pandas as pd
from tqdm.notebook import tqdm
```

```
# enable tqdm in pandas
tqdm.pandas()
```

```
# set to True to use the gpu (if there is one available)
use_gpu = True
```

```
# select device
device = torch.device('cuda' if use_gpu and torch.cuda.is_available() else 'cpu')
print(f'device: {device.type}')
```

```
# random seed
seed = 1234
```

```
# set random seed
if seed is not None:
    print(f'random seed: {seed}')
    random.seed(seed)
    np.random.seed(seed)
    torch.manual_seed(seed)
```

```
↗ device: cpu
random seed: 1234
```

We will be using the AG's News Topic Classification Dataset. It is stored in two CSV files: `train.csv` and `test.csv`, as well as a `classes.txt` that stores the labels of the classes to predict.

First, we will load the training dataset using [pandas](#) and take a quick look at how the data.

```
train_df = pd.read_csv('train.csv', header=None)
train_df.columns = ['class index', 'title', 'description']
train_df
```



| | class index | | title | description |
|--------|-------------|---|---|-------------|
| 0 | 3 | Wall St. Bears Claw Back Into the Black (Reuters) | Reuters - Short-sellers, Wall Street's dwindli... | |
| 1 | 3 | Carlyle Looks Toward Commercial Aerospace (Reu... | Reuters - Private investment firm Carlyle Grou... | |
| 2 | 3 | Oil and Economy Cloud Stocks' Outlook (Reuters) | Reuters - Soaring crude prices plus worries\ab... | |
| 3 | 3 | Iraq Halts Oil Exports from Main Southern Pipe... | Reuters - Authorities have halted oil exportf... | |
| 4 | 3 | Oil prices soar to all-time record, posing new... | AFP - Tearaway world oil prices, toppling reco... | |
| ... | ... | | | ... |
| 119995 | 1 | Pakistan's Musharraf Says Won't Quit as Army C... | KARACHI (Reuters) - Pakistani President Perve... | |
| 119996 | 2 | Renteria signing a top-shelf deal | Red Sox general manager Theo Epstein acknowle... | |
| 119997 | 2 | Saban not going to Dolphins yet | The Miami Dolphins will put their courtship of... | |
| 119998 | 2 | Today's NFL games | PITTSBURGH at NY GIANTS Time: 1:30 p.m. Line: ... | |
| 119999 | 2 | Nets get Carter from Raptors | INDIANAPOLIS -- All-Star Vince Carter was trad... | |

120000 rows x 3 columns

The dataset consists of 120,000 examples, each consisting of a class index, a title, and a description. The class labels are distributed in a separated file. We will add the labels to the dataset so that we can interpret the data more easily. Note that the label indexes are one-based, so we need to subtract one to retrieve them from the list.

```
labels = open('classes.txt').read().splitlines()
classes = train_df['class index'].map(lambda i: labels[i-1])
train_df.insert(1, 'class', classes)
train_df
```



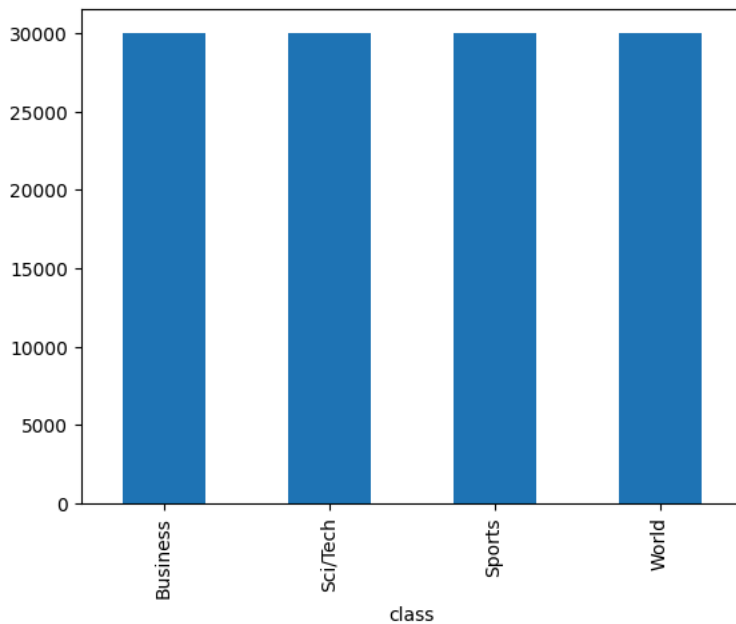
| | class index | class | | title | description |
|--------|-------------|----------|---|---|-------------|
| 0 | 3 | Business | Wall St. Bears Claw Back Into the Black (Reuters) | Reuters - Short-sellers, Wall Street's dwindli... | |
| 1 | 3 | Business | Carlyle Looks Toward Commercial Aerospace (Reu... | Reuters - Private investment firm Carlyle Grou... | |
| 2 | 3 | Business | Oil and Economy Cloud Stocks' Outlook (Reuters) | Reuters - Soaring crude prices plus worries\ab... | |
| 3 | 3 | Business | Iraq Halts Oil Exports from Main Southern Pipe... | Reuters - Authorities have halted oil exportf... | |
| 4 | 3 | Business | Oil prices soar to all-time record, posing new... | AFP - Tearaway world oil prices, toppling reco... | |
| ... | ... | ... | | ... | ... |
| 119995 | 1 | World | Pakistan's Musharraf Says Won't Quit as Army C... | KARACHI (Reuters) - Pakistani President Perve... | |
| 119996 | 2 | Sports | Renteria signing a top-shelf deal | Red Sox general manager Theo Epstein acknowle... | |
| 119997 | 2 | Sports | Saban not going to Dolphins yet | The Miami Dolphins will put their courtship of... | |
| 119998 | 2 | Sports | Today's NFL games | PITTSBURGH at NY GIANTS Time: 1:30 p.m. Line: ... | |
| 119999 | 2 | Sports | Nets get Carter from Raptors | INDIANAPOLIS -- All-Star Vince Carter was trad... | |

120000 rows x 4 columns

Let's inspect how balanced our examples are by using a bar plot.

```
pd.value_counts(train_df['class']).plot.bar()
```

```
<ipython-input-5-78d1f3cd1886>:1: FutureWarning: pandas.value_counts is deprecated and will be removed in a future version.
pd.value_counts(train_df['class']).plot.bar()
<Axes: xlabel='class'>
```



The classes are evenly distributed. That's great!

However, the text contains some spurious backslashes in some parts of the text. They are meant to represent newlines in the original text. An example can be seen below, between the words "dwindling" and "band".

```
print(train_df.loc[0, 'description'])
```

```
Reuters - Short-sellers, Wall Street's dwindling\band of ultra-cynics, are seeing green again.
```

We will replace the backslashes with spaces on the whole column using pandas replace method.

```
title = train_df['title'].str.lower()
descr = train_df['description'].str.lower()
text = title + " " + descr
train_df['text'] = text.str.replace('\\', ' ', regex=False)
train_df
```

| | class index | class | title | description | text |
|--------|-------------|----------|---|---|---|
| 0 | 3 | Business | Wall St. Bears Claw Back Into the Black (Reuters) | Reuters - Short-sellers, Wall Street's dwindli... | wall st. bears claw back into the black (reute... |
| 1 | 3 | Business | Carlyle Looks Toward Commercial Aerospace (Reu... | Reuters - Private investment firm Carlyle Grou... | carlyle looks toward commercial aerospace (reu... |
| 2 | 3 | Business | Oil and Economy Cloud Stocks' Outlook (Reuters) | Reuters - Soaring crude prices plus worries\ab... | oil and economy cloud stocks' outlook (reuters... |
| 3 | 3 | Business | Iraq Halts Oil Exports from Main Southern Pipe... | Reuters - Authorities have halted oil exportf... | iraq halts oil exports from main southern pipe... |
| 4 | 3 | Business | Oil prices soar to all-time record, posing new... | AFP - Tearaway world oil prices, toppling reco... | oil prices soar to all-time record, posing new... |
| ... | ... | ... | ... | ... | ... |
| 119995 | 1 | World | Pakistan's Musharraf Says Won't Quit as Army C... | KARACHI (Reuters) - Pakistani President Perve... | pakistan's musharraf says won't quit as army c... |
| 119996 | 2 | Sports | Renteria signing a top-shelf deal | Red Sox general manager Theo Epstein acknowle... | renteria signing a top-shelf deal red sox gene... |

Now we will proceed to tokenize the title and description columns using NLTK's word_tokenize(). We will add a new column to our dataframe with the list of tokens.

```
import nltk
nltk.download('punkt')
```

```
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Package punkt is already up-to-date!
True
```

```
from nltk.tokenize import word_tokenize
```

```
train_df['tokens'] = train_df['text'].progress_map(word_tokenize)
train_df
```

```
100% 120000/120000 [00:44<00:00, 3221.73it/s]
```

| | class index | class | title | description | text | tokens |
|--------|-------------|----------|---|---|---|---|
| 0 | 3 | Business | Wall St. Bears Claw Back Into the Black (Reuters) | Reuters - Short-sellers, Wall Street's dwindli... | wall st. bears claw back into the black (reute... | [wall, st., bears, claw, back, into, the, blac... |
| 1 | 3 | Business | Carlyle Looks Toward Commercial Aerospace (Reu... | Reuters - Private investment firm Carlyle Grou... | carlyle looks toward commercial aerospace (reu... | [carlyle, looks, toward, commercial, aerospace... |
| 2 | 3 | Business | Oil and Economy Cloud Stocks' Outlook (Reuters) | Reuters - Soaring crude prices plus worries\ab... | oil and economy cloud stocks' outlook (reuters... | [oil, and, economy, cloud, stocks, ', outlook,... |
| 3 | 3 | Business | Iraq Halts Oil Exports from Main Southern Pipe... | Reuters - Authorities have halted oil exportf... | iraq halts oil exports from main southern pipe... | [iraq, halts, oil, exports, from, main, southe... |
| 4 | 3 | Business | Oil prices soar to all-time record, posing new... | AFP - Tearaway world oil prices, toppling reco... | oil prices soar to all-time record, posing new... | [oil, prices, soar, to, all-time, record, ,, p... |
| ... | ... | ... | ... | ... | ... | ... |
| 119995 | 1 | World | Pakistan's Musharraf Says Won't Quit as Army C... | KARACHI (Reuters) - Pakistani President Perve... | pakistan's musharraf says won't quit as army c... | [pakistan, 's, musharraf, says, wo, n't, quit,... |
| 119996 | 2 | Sports | Renteria signing a top-shelf deal | Red Sox general manager Theo Epstein acknowled... | renteria signing a top-shelf deal red sox gene... | [renteria, signing, a, top-shelf, deal, red, s... |
| 119997 | 2 | Sports | Saban not going to Dolphins yet | The Miami Dolphins will put their courtship of... | saban not going to dolphins yet the miami dorp... | [saban, not, going, to, dolphins, yet, the, mi... |
| 119998 | 2 | Sports | Today's NFL games | PITTSBURGH at NY GIANTS Time: 1:30 p.m. Line: ... | today's nfl games pittsburgh at ny giants time... | [today, 's, nfl, games, pittsburgh, at, ny, gi... |
| ---- | - | - | - | INDIANAPOLIS -- All-Star Vince | nets net carter from rantors | [nets, net, carter, from |

Now we will create a vocabulary from the training data. We will only keep the terms that repeat beyond some threshold established below.

```
threshold = 10
tokens = train_df['tokens'].explode().value_counts()
tokens = tokens[tokens > threshold]
id_to_token = ['[UNK]'] + tokens.index.tolist()
token_to_id = {w:i for i,w in enumerate(id_to_token)}
vocabulary_size = len(id_to_token)
print(f'vocabulary size: {vocabulary_size:,}')
```

```
vocabulary size: 19,671
```

```
from collections import defaultdict
```

```
def make_feature_vector(tokens, unk_id=0):
    vector = defaultdict(int)
    for t in tokens:
        i = token_to_id.get(t, unk_id)
        vector[i] += 1
    return vector
```

```
train_df['features'] = train_df['tokens'].progress_map(make_feature_vector)
train_df
```



100%

120000/120000 [00:05<00:00, 23774.27it/s]

| | class index | class | title | description | text | tokens | features |
|--------|-------------|----------|---|---|---|---|--|
| 0 | 3 | Business | Wall St. Bears Claw Back Into the Black (Reuters) | Reuters - Short-sellers, Wall Street's dwindli... | wall st. bears claw back into the black (reute... | [wall, st., bears, claw, back, into, the, blac... | {427: 2, 563: 1, 1607: 1, 15337: 1, 120: 1, 73... |
| 1 | 3 | Business | Carlyle Looks Toward Commercial Aerospace (Reu... | Reuters - Private investment firm Carlyle Grou... | carlyle looks toward commercial aerospace (reu... | [carlyle, looks, toward, commercial, aerospace... | {16358: 2, 1078: 1, 855: 1, 1287: 1, 4248: 1, ... |
| 2 | 3 | Business | Oil and Economy Cloud Stocks' Outlook (Reuters) | Reuters - Soaring crude prices plus worrieslab... | oil and economy cloud stocks' outlook (reuters... | [oil, and, economy, cloud, stocks, ', outlook,... | {66: 1, 9: 2, 351: 2, 4564: 1, 158: 1, 116: 1,... |
| 3 | 3 | Business | Iraq Halts Oil Exports from Main Southern Pipe... | Reuters - Authorities have halted oil exportf... | iraq halts oil exports from main southern pipe... | [iraq, halts, oil, exports, from, main, southe... | {77: 2, 7372: 1, 66: 3, 1783: 1, 32: 2, 900: 2... |
| 4 | 3 | Business | Oil prices soar to all-time record, posing new... | AFP - Tearaway world oil prices, toppling reco... | oil prices soar to all-time record, posing new... | [oil, prices, soar, to, all-time, record, ,, p... | {66: 2, 99: 2, 4387: 1, 4: 2, 3604: 1, 149: 1,... |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 119995 | 1 | World | Pakistan's Musharraf Says Won't Quit as Army C... | KARACHI (Reuters) - Pakistani President Perve... | pakistan's musharraf says won't quit as army c... | [pakistan, 's, musharraf, says, wo, n't, quit,... | {383: 1, 23: 1, 1625: 2, 91: 1, 1804: 1, 285: ... |
| 119996 | 2 | Sports | Renteria signing a top-shelf deal | Red Sox general manager Theo Epstein acknowled... | renteria signing a top-shelf deal red sox gene... | [renteria, signing, a, top-shelf, deal, red, s... | {8510: 2, 2637: 1, 5: 4, 0: 3, 127: 1, 202: 3,... |
| 119997 | 2 | Sports | Saban not going to Dolphins yet | The Miami Dolphins will put their courtship of... | saban not going to dolphins yet the miami doli... | [saban, not, going, to, dolphins, yet, the, mi... | {7758: 2, 68: 1, 661: 1, 4: 2, 1439: 2, 704: 1,... |

Se tokeniza el texto con word_tokenize de NLTK y se agrega como una nueva columna al DataFrame. Se crea un vocabulario a partir de los términos que aparecen con una frecuencia mayor a un umbral de 10. Cada texto se convierte en un vector de características donde cada posición representa la frecuencia de un token específico en el vocabulario.

```
def make_dense(feats):
    x = np.zeros(vocabulary_size)
    for k,v in feats.items():
        x[k] = v
    return x

X_train = np.stack(train_df['features'].progress_map(make_dense))
y_train = train_df['class index'].to_numpy() - 1

X_train = torch.tensor(X_train, dtype=torch.float32)
y_train = torch.tensor(y_train)
```



66%

79253/120000 [00:10<00:08, 4758.11it/s]

```
from torch import nn
from torch import optim

# hyperparameters
lr = 1.0
n_epochs = 5
n_examples = X_train.shape[0]
n_feats = X_train.shape[1]
n_classes = len(labels)

# initialize the model, loss function, optimizer, and data-loader
model = nn.Linear(n_feats, n_classes).to(device)
loss_func = nn.CrossEntropyLoss()
optimizer = optim.SGD(model.parameters(), lr=lr)

# train the model
indices = np.arange(n_examples)
for epoch in range(n_epochs):
    np.random.shuffle(indices)
    for i in tqdm(indices, desc=f'epoch {epoch+1}'):
        # clear gradients
        # ...
```

```

model.zero_grad()
# send datum to right device
x = X_train[i].unsqueeze(0).to(device)
y_true = y_train[i].unsqueeze(0).to(device)
# predict label scores
y_pred = model(x)
# compute loss
loss = loss_func(y_pred, y_true)
# backpropagate
loss.backward()
# optimize model parameters
optimizer.step()

```

```

epoch 1: 0%|          | 0/120000 [00:00<?, ?it/s]
epoch 2: 0%|          | 0/120000 [00:00<?, ?it/s]
epoch 3: 0%|          | 0/120000 [00:00<?, ?it/s]
epoch 4: 0%|          | 0/120000 [00:00<?, ?it/s]
epoch 5: 0%|          | 0/120000 [00:00<?, ?it/s]

```

Se convierte la matriz de características y las etiquetas a tensores de PyTorch para su uso en el modelo. Se define un modelo de regresión logística con PyTorch, que incluye una función de pérdida de entropía cruzada (CrossEntropyLoss) y el optimizador SGD. Luego se entrena el modelo en lotes, realizando predicciones, calculando el error, y ajustando los parámetros.

Next, we evaluate on the test dataset

Se aplica el mismo flujo de preprocesamiento de texto al conjunto de prueba y se convierte en tensores. Se utiliza `classification_report` de Scikit-learn para evaluar el rendimiento del modelo, generando métricas de clasificación, como precisión, recall, y F1, para cada clase.

[+ Código](#)
[+ Texto](#)

```

# repeat all preprocessing done above, this time on the test set
test_df = pd.read_csv('data/ag_news_csv/test.csv', header=None)
test_df.columns = ['class index', 'title', 'description']
test_df['text'] = test_df['title'].str.lower() + " " + test_df['description'].str.lower()
test_df['text'] = test_df['text'].str.replace('\n', ' ', regex=False)
test_df['tokens'] = test_df['text'].progress_map(word_tokenize)
test_df['features'] = test_df['tokens'].progress_map(make_feature_vector)

X_test = np.stack(test_df['features'].progress_map(make_dense))
y_test = test_df['class index'].to_numpy() - 1
X_test = torch.tensor(X_test, dtype=torch.float32)
y_test = torch.tensor(y_test)

```

```

epoch 1: 0%|          | 0/7600 [00:00<?, ?it/s]
epoch 2: 0%|          | 0/7600 [00:00<?, ?it/s]
epoch 3: 0%|          | 0/7600 [00:00<?, ?it/s]

```

```
from sklearn.metrics import classification_report
```

```
# set model to evaluation mode
model.eval()
```

```

# don't store gradients
with torch.no_grad():
    X_test = X_test.to(device)
    y_pred = torch.argmax(model(X_test), dim=1)
    y_pred = y_pred.cpu().numpy()
    print(classification_report(y_test, y_pred, target_names=labels))

```

```

precision    recall  f1-score   support

World        0.94      0.82      0.88      1900
Sports       0.89      0.99      0.94      1900
Business     0.81      0.88      0.85      1900
Sci/Tech     0.89      0.83      0.86      1900

```